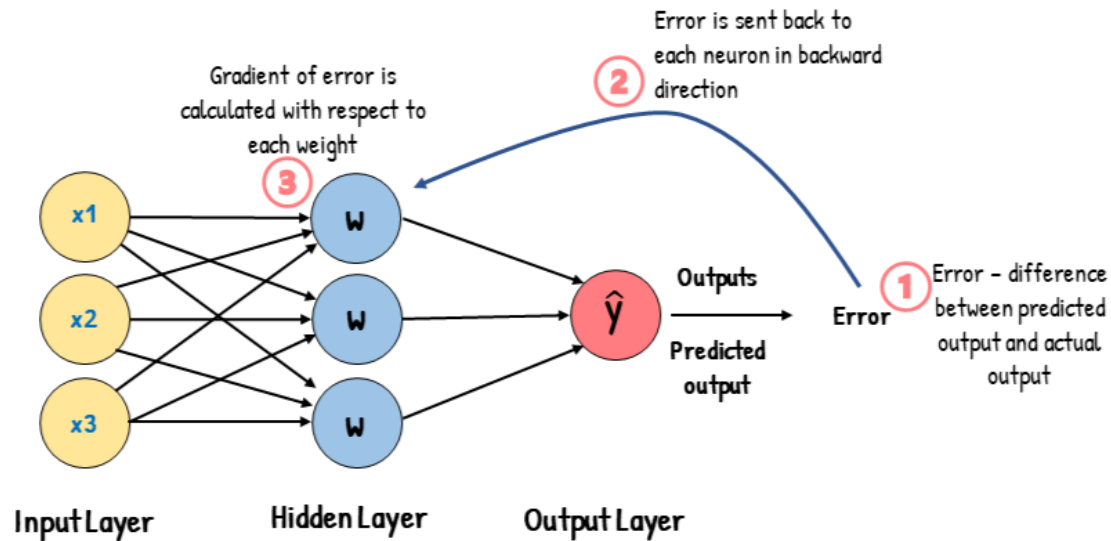


Backpropagation



Introduction to Optimisation for ML

Derivatives and Gradients


Andreas Makris, 2nd year PhD student
Lancaster University, ProbAI Hub

Based on chapter 2 of the book **Algorithms for Optimization** by Mykel J. Kochenderfer Tim A. Wheeler

What is optimization?


- We have a function f that depends on some input x . We want to find x that minimizes f subject to some constrain. Mathematically:

$$\begin{array}{ll}\underset{\mathbf{x}}{\text{minimize}} & f(\mathbf{x}) \\ \text{subject to} & \mathbf{x} \in \mathcal{X}\end{array}$$



Why do we need optimization in ML?


- The function is the loss function J . The input are the parameters of the model θ .



Why do we need optimization in ML?

- The function is the loss function J . The input are the parameters of the model θ .
- We want to find the parameters of the model that minimize the loss function.

GRADIENT DESCENT ALGORITHM



$$\theta := \theta - \alpha \nabla J(\theta)$$

θ \longrightarrow parameter

α \longrightarrow learning rate


$\nabla J(\theta)$ \longrightarrow gradient of cost function

Why do we need optimization in ML?

- The function is the loss function J . The input are the parameters of the model θ .
- We want to find the parameters of the model that minimize the loss function.
- There are a lot of optimization algorithms that use the gradient of the function with respect to the input (e.g. gradient descent, ADAM).
- Today we will focus on how to calculate the gradient of the loss with respect to the parameters of the model.

$$\frac{\partial J}{\partial \theta}$$

GRADIENT DESCENT ALGORITHM



$$\theta := \theta - \alpha \nabla J(\theta)$$

θ \longrightarrow parameter


α \longrightarrow learning rate

$\nabla J(\theta)$ \longrightarrow gradient of cost function

Automatic Differentiation

- Two types (modes) of autodiff; forward mode and reverse mode (backpropagation).

$$f(a, b) = \ln(ab + \max(a, 2))$$




$$c_4 = \ln(c_3) = \ln(ab + \max(a, 2)) \quad c_3 = c_1 + c_2 = ab + \max(a, 2)$$

Automatic Differentiation

- Two types (modes) of autodiff; forward mode and reverse mode (backpropagation).
- Can be used when a function can be expressed as a computation graph with all elementary functions being differentiable.

$$f(a, b) = \ln(ab + \max(a, 2))$$




$$c_4 = \ln(c_3) = \ln(ab + \max(a, 2)) \quad c_3 = c_1 + c_2 = ab + \max(a, 2)$$

Automatic Differentiation

- Two types (modes) of autodiff; forward mode and reverse mode (backpropagation).
- Can be used when a function can be expressed as a computation graph with all elementary functions being differentiable.
- Start by building the computation graph; inputs on the left, operations are nodes, introduce intermediate variables, outputs on the right.

$$f(a, b) = \ln(ab + \max(a, 2))$$




$$c_4 = \ln(c_3) = \ln(ab + \max(a, 2)) \quad c_3 = c_1 + c_2 = ab + \max(a, 2)$$

Automatic Differentiation

- Two types (modes) of autodiff; forward mode and reverse mode (backpropagation).
- Can be used when a function can be expressed as a computation graph with all elementary functions being differentiable.
- Start by building the computation graph; inputs on the left, operations are nodes, introduce intermediate variables, outputs on the right.
- Both modes are based on the **chain rule**.
- Our goal is to calculate $\frac{\partial f}{\partial \alpha}$ (and $\frac{\partial f}{\partial b}$).

$$f(a, b) = \ln(ab + \max(a, 2))$$




$$c_4 = \ln(c_3) = \ln(ab + \max(a, 2)) \quad c_3 = c_1 + c_2 = ab + \max(a, 2)$$

Automatic Differentiation

- Forward mode: Calculate in order $\frac{\partial c_1}{\partial \alpha}$, $\frac{\partial c_2}{\partial \alpha}$, $\frac{\partial c_3}{\partial \alpha}$, $\frac{\partial f}{\partial \alpha}$.

$$f(a, b) = \ln(ab + \max(a, 2))$$




$$c_4 = \ln(c_3) = \ln(ab + \max(a, 2)) \quad c_3 = c_1 + c_2 = ab + \max(a, 2)$$

Automatic Differentiation

- Forward mode: Calculate in order $\frac{\partial c_1}{\partial \alpha}$, $\frac{\partial c_2}{\partial \alpha}$, $\frac{\partial c_3}{\partial \alpha}$, $\frac{\partial f}{\partial \alpha}$.
- Reverse mode: Calculate in order $\frac{\partial f}{\partial c_3}$, $\frac{\partial f}{\partial c_2}$, $\frac{\partial f}{\partial c_1}$, $\frac{\partial f}{\partial \alpha}$.

$$f(a, b) = \ln(ab + \max(a, 2))$$




$$c_4 = \ln(c_3) = \ln(c_1 + c_2) = \ln(ab + \max(a, 2)) \quad c_3 = c_1 + c_2$$

Automatic Differentiation

- Forward mode: Calculate in order $\frac{\partial c_1}{\partial \alpha}$, $\frac{\partial c_2}{\partial \alpha}$, $\frac{\partial c_3}{\partial \alpha}$, $\frac{\partial f}{\partial \alpha}$.
- Reverse mode: Calculate in order $\frac{\partial f}{\partial c_3}$, $\frac{\partial f}{\partial c_2}$, $\frac{\partial f}{\partial c_1}$, $\frac{\partial f}{\partial \alpha}$.
- When the input dimensionality is higher than the output dimensionality reverse mode is cheaper.
- When the input dimensionality is lower than the output dimensionality forward mode is cheaper.

$$f(a, b) = \ln(ab + \max(a, 2))$$




$$c_4 = \ln(c_3) = \ln(ab + \max(a, 2)) \quad c_3 = c_1 + c_2 = ab + \max(a, 2)$$

Forward Mode

- Let $a=3$ and $b=2$. Use the forward mode autodiff to find $\frac{\partial f}{\partial a}$.

$$f(a, b) = \ln(ab + \max(a, 2))$$



Forward Mode

- Let $a=3$ and $b=2$. Use the forward mode autodiff to find $\frac{\partial f}{\partial \alpha}$.
- For each node calculate both the value and partial derivative with respect to a .


- Example use of chain rule;

$$\frac{\partial c_3}{\partial \alpha} = \frac{\partial c_3}{\partial c_1} \frac{\partial c_1}{\partial \alpha} + \frac{\partial c_3}{\partial c_2} \frac{\partial c_2}{\partial \alpha}$$

$$c_4 = \ln(c_3) = \ln(c_1 + c_2) = \ln(ab + \max(a, 2)) \quad c_3 = c_1 + c_2 = ab + \max(a, 2) \quad c_2 = \max(a, 2) \quad c_1 = ab$$

$$\dot{c}_2 = \begin{cases} 0 & \text{if } 2 > a \\ \dot{a} & \text{if } 2 < a \end{cases} = 1$$

$$f(a, b) = \ln(ab + \max(a, 2))$$




Reverse Mode

- Start by a forward pass to calculate the values (only).
- Do a reverse pass for the gradients.

$$\begin{aligned}\frac{\partial f}{\partial c} &= \frac{\partial f}{\partial d} \frac{\partial d}{\partial c} + \frac{\partial f}{\partial e} \frac{\partial e}{\partial c} \\ \frac{\partial f}{\partial b} &= \frac{\partial f}{\partial c} \frac{\partial c}{\partial b} \\ \frac{\partial f}{\partial a} &= \frac{\partial f}{\partial b} \frac{\partial b}{\partial a} + \frac{\partial f}{\partial c} \frac{\partial c}{\partial a} \\ \frac{\partial f}{\partial x} &= \frac{\partial f}{\partial a} \frac{\partial a}{\partial x}.\end{aligned}$$

$$f(x) = \sqrt{x^2 + \exp(x^2)} + \cos(x^2 + \exp(x^2))$$




- This is what neural networks use to calculate the gradients.

Reverse Mode

- Start by a forward pass to calculate the values.
- Do a reverse pass for the gradients.

$$\begin{aligned}\frac{\partial f}{\partial c} &= \frac{\partial f}{\partial d} \frac{\partial d}{\partial c} + \frac{\partial f}{\partial e} \frac{\partial e}{\partial c} \\ \frac{\partial f}{\partial b} &= \frac{\partial f}{\partial c} \frac{\partial c}{\partial b} \\ \frac{\partial f}{\partial a} &= \frac{\partial f}{\partial b} \frac{\partial b}{\partial a} + \frac{\partial f}{\partial c} \frac{\partial c}{\partial a} \\ \frac{\partial f}{\partial x} &= \frac{\partial f}{\partial a} \frac{\partial a}{\partial x}.\end{aligned}$$

$$f(x) = \sqrt{x^2 + \exp(x^2)} + \cos(x^2 + \exp(x^2))$$



- This is what neural networks use to calculate the gradients.

What if f cannot be expressed as a computational graph with differentiable functions?

Numerical Differentiation

- Estimate derivatives numerically (not exact!!).
- **Finite difference methods;** use the definition of differentiation and plug in a small value of h .
- Forward difference $O(h)$ but central difference $O(h^2)$.
- If h is too small, we might face numerical subtractive cancellation issues.

$$f'(x) \approx \underbrace{\frac{f(x+h) - f(x)}{h}}_{\text{forward difference}} \approx \underbrace{\frac{f(x+h/2) - f(x-h/2)}{h}}_{\text{central difference}} \approx \underbrace{\frac{f(x) - f(x-h)}{h}}_{\text{backward difference}}$$

Numerical Differentiation

- Estimate derivatives numerically (not exact!!).
- **Finite difference methods;** use the definition of differentiation and plug in a small value of h .
- Forward difference $O(h)$ but central difference $O(h^2)$.
- If h is too small, we might face numerical subtractive cancellation issues.
- **Complex step method;**

$$f'(x) \approx \frac{\text{Im}(f(x + ih))}{h}$$

- No subtractive cancellation issues. $O(h^2)$.
- All proofs with Taylor series.

$$f'(x) \approx \underbrace{\frac{f(x+h) - f(x)}{h}}_{\text{forward difference}} \approx \underbrace{\frac{f(x+h/2) - f(x-h/2)}{h}}_{\text{central difference}} \approx \underbrace{\frac{f(x) - f(x-h)}{h}}_{\text{backward difference}}$$

Numerical Differentiation

- Estimate derivatives numerically (not exact!!).
- **Finite difference methods;** use the definition of differentiation and plug in a small value of h .
- Forward difference $O(h)$ but central difference $O(h^2)$.
- If h is too small, we might face numerical subtractive cancellation issues.
- **Complex step method;**

$$f'(x) \approx \frac{\text{Im}(f(x + ih))}{h}$$

- No subtractive cancellation issues. $O(h^2)$.
- All proofs with Taylor series.

$$f'(x) \approx \underbrace{\frac{f(x+h) - f(x)}{h}}_{\text{forward difference}} \approx \underbrace{\frac{f(x+h/2) - f(x-h/2)}{h}}_{\text{central difference}} \approx \underbrace{\frac{f(x) - f(x-h)}{h}}_{\text{backward difference}}$$

These methods do not scale well with the number of parameters.

Regression Gradient

- Used for problems with noisy objective functions because the regression helps smooth out the noise when producing a gradient estimate.
- Need a dataset of perturbations and their function evaluations.
- Use first-order Taylor expansion.
- Find \mathbf{g} using linear regression.

$$\mathbf{g} = \Delta \mathbf{X}^+ \Delta \mathbf{f}$$

$$\left(\Delta \mathbf{x}^{(1)}, f(\mathbf{x} + \Delta \mathbf{x}^{(1)}) \right), \left(\Delta \mathbf{x}^{(2)}, f(\mathbf{x} + \Delta \mathbf{x}^{(2)}) \right), \dots, \left(\Delta \mathbf{x}^{(m)}, f(\mathbf{x} + \Delta \mathbf{x}^{(m)}) \right)$$

$$\hat{f}(\mathbf{x} + \Delta \mathbf{x}) = f(\mathbf{x}) + \mathbf{g}^\top \Delta \mathbf{x}$$

$$\Delta \mathbf{X} = \begin{bmatrix} (\Delta \mathbf{x}^{(1)})^\top \\ \vdots \\ (\Delta \mathbf{x}^{(m)})^\top \end{bmatrix}$$

$$\Delta \mathbf{f} = \left[f(\mathbf{x} + \Delta \mathbf{x}^{(1)}) - f(\mathbf{x}), \dots, f(\mathbf{x} + \Delta \mathbf{x}^{(m)}) - f(\mathbf{x}) \right]$$

Directional Derivatives

- “The directional derivative $\nabla_s f(\mathbf{x})$ of a multivariate function f is the instantaneous rate of change of $f(\mathbf{x})$ as \mathbf{x} is moved with velocity \mathbf{s} .”
- We can calculate the directional derivative using the following formula.

$$\nabla_s f(\mathbf{x}) = \nabla f(\mathbf{x})^\top \mathbf{s}$$

- The directional derivative is a scalar (when the function output is scalar)!
- The directional derivative is highest in the gradient direction and lowest in the opposite direction of the gradient.

$$\nabla_s f(\mathbf{x}) \equiv \underbrace{\lim_{h \rightarrow 0} \frac{f(\mathbf{x} + h\mathbf{s}) - f(\mathbf{x})}{h}}_{\text{forward difference}} = \underbrace{\lim_{h \rightarrow 0} \frac{f(\mathbf{x} + h\mathbf{s}/2) - f(\mathbf{x} - h\mathbf{s}/2)}{h}}_{\text{central difference}} = \underbrace{\lim_{h \rightarrow 0} \frac{f(\mathbf{x}) - f(\mathbf{x} - h\mathbf{s})}{h}}_{\text{backward difference}}$$

Simultaneous Perturbation Stochastic Gradient Approximation (SPSA)

- SPSA can estimate the gradient with as few as two function evaluations, regardless of the number of variables. Can work in deep learning.
- Uses directional derivatives.

$$\nabla_{\mathbf{z}} f(\mathbf{x}) \approx \frac{f(\mathbf{x} + \delta \mathbf{z}) - f(\mathbf{x} - \delta \mathbf{z})}{2\delta}$$

$$\nabla f(\mathbf{x}) \approx (\nabla_{\mathbf{z}} f(\mathbf{x})) \mathbf{z}$$

Simultaneous Perturbation Stochastic Gradient Approximation (SPSA)

- SPSA can estimate the gradient with as few as two function evaluations, regardless of the number of variables. Can work in deep learning.
- Uses directional derivatives.
- $\mathbf{z} \sim N(0, I)$.
- Average many samples to improve the estimate.
- “The sample count is typically left quite small or even set to 1”.

$$\nabla_{\mathbf{z}} f(\mathbf{x}) \approx \frac{f(\mathbf{x} + \delta \mathbf{z}) - f(\mathbf{x} - \delta \mathbf{z})}{2\delta}$$

$$\nabla f(\mathbf{x}) \approx (\nabla_{\mathbf{z}} f(\mathbf{x})) \mathbf{z}$$

Thank you for
listening!
Questions?

